

Deep Network Approximation with Discrepancy Being Reciprocal of Width to Power of Depth

Zuowei Shen

matzuows@nus.edu.sg

Department of Mathematics, National University of Singapore

Haizhao Yang

haizhao@purdue.edu

Department of Mathematics, Purdue University

Shijun Zhang

zhangshijun@u.nus.edu

Department of Mathematics, National University of Singapore

Keywords: Exponential Convergence, Curse of Dimensionality, Deep Neural Network, Floor-ReLU Activation Function, Approximation via Compositions, Continuous Function.

Abstract

A new network with super approximation power is introduced. This network is built with Floor ($\lfloor x \rfloor$) and ReLU ($\max\{0, x\}$) activation functions and hence we call such networks as Floor-ReLU networks. It is shown by construction that Floor-ReLU networks with width $\max\{d, 5N + 13\}$ and depth $64dL + 3$ can pointwise approximate a Lipschitz continuous function f on $[0, 1]^d$ with an exponential approximation rate $3\mu\sqrt{d}N^{-\sqrt{L}}$, where μ is the Lipschitz constant of f . More generally for an arbitrary continuous function f on $[0, 1]^d$ with a modulus of continuity $\omega_f(\cdot)$, the constructive approximation rate is $\omega_f(\sqrt{d}N^{-\sqrt{L}}) + 2\omega_f(\sqrt{d})N^{-\sqrt{L}}$. As a consequence, this new network overcomes the curse of dimensionality in approximation power since this approximation order is essentially \sqrt{d} times a function of N and L independent of d .

1 Introduction

Recently, there has been a large number of successful real-world applications of deep neural networks in many fields of computer science and engineering, especially for

large-scale and high-dimensional learning problems. Understanding the approximation capacity of deep neural networks has become a fundamental research direction for revealing the advantages of deep learning versus traditional methods. This paper introduces new theories and network architectures achieving exponential convergence and avoiding the curse of dimensionality for continuous functions for the first time in deep network approximation, which might be two foundational laws supporting the application of deep network approximation in large-scale and high-dimensional problems. The approximation theories here are quantitative and work for networks with essentially arbitrary width and depth. They would shed new light on the design of the efficient numerical implementation of deep learning.

Deep ReLU networks can achieve the approximation rate $O(N^{-L})$ for polynomials on $[0, 1]^d$ (Lu et al., 2020) but it is not true for general functions, e.g., the optimal approximation rates of deep ReLU networks for continuous functions and C^s functions f on $[0, 1]^d$ are $O(\sqrt{d}N^{-2/d}L^{-2/d})$ and $O(\|f\|_{C^s}N^{-2s/d}L^{-2s/d})$ (Shen et al., 2019; Lu et al., 2020), respectively. The limitation of ReLU networks motivates us to explore other types of network architectures to seek the answers to two fundamental questions: Do deep neural networks with an arbitrary width N and an arbitrary depth L admit an approximation rate $O(c(d)N^{-L})$ for general functions in a d -dimensional space? How small the function $c(d)$ in d could be?

In particular, we introduce the Floor-ReLU network, which is a fully connected neural network (FNN) built with either Floor ($\lfloor x \rfloor$) or ReLU ($\max\{0, x\}$) activation function¹ in each neuron. Mathematically, if we let $N_0 = d$, $N_{L+1} = 1$, and N_ℓ be the number of neurons in ℓ -th hidden layer of a Floor-ReLU network for $\ell = 1, 2, \dots, L$, then the architecture of this network with input \mathbf{x} and output $\phi(\mathbf{x})$ can be described as

$$\mathbf{x} = \tilde{\mathbf{h}}_0 \xrightarrow{\mathbf{W}_0, \mathbf{b}_0} \mathbf{h}_1 \xrightarrow{\sigma \text{ or } \lfloor \cdot \rfloor} \tilde{\mathbf{h}}_1 \quad \dots \quad \xrightarrow{\mathbf{W}_{L-1}, \mathbf{b}_{L-1}} \mathbf{h}_L \xrightarrow{\sigma \text{ or } \lfloor \cdot \rfloor} \tilde{\mathbf{h}}_L \xrightarrow{\mathbf{W}_L, \mathbf{b}_L} \mathbf{h}_{L+1} = \phi(\mathbf{x}),$$

where $\mathbf{W}_\ell \in \mathbb{R}^{N_{\ell+1} \times N_\ell}$, $\mathbf{b}_\ell \in \mathbb{R}^{N_{\ell+1}}$, $\mathbf{h}_{\ell+1} := \mathbf{W}_\ell \cdot \tilde{\mathbf{h}}_\ell + \mathbf{b}_\ell$ for $\ell = 0, 1, \dots, L$, and $\tilde{\mathbf{h}}_{\ell,n}$ equals to $\sigma(\mathbf{h}_{\ell,n})$ or $\lfloor \mathbf{h}_{\ell,n} \rfloor$ for $\ell = 1, 2, \dots, L$ and $n = 1, 2, \dots, N_\ell$, where $\mathbf{h}_\ell = (\mathbf{h}_{\ell,1}, \dots, \mathbf{h}_{\ell,N_\ell})$ and $\tilde{\mathbf{h}}_\ell = (\tilde{\mathbf{h}}_{\ell,1}, \dots, \tilde{\mathbf{h}}_{\ell,N_\ell})$ for $\ell = 1, 2, \dots, L$.

In Theorem 1.1 below, we show by construction that Floor-ReLU networks with width $\max\{d, 5N+13\}$ and depth $64dL+3$ can pointwise approximate an arbitrary continuous function f on $[0, 1]^d$ with an exponential approximation rate $\omega_f(\sqrt{d}N^{-\sqrt{L}}) + 2\omega_f(\sqrt{d})N^{-\sqrt{L}}$, where $\omega_f(\cdot)$ is the modulus of continuity defined as

$$\omega_f(r) := \sup \{ |f(\mathbf{x}) - f(\mathbf{y})| : \|\mathbf{x} - \mathbf{y}\|_2 \leq r, \mathbf{x}, \mathbf{y} \in [0, 1]^d \}, \quad \text{for any } r \geq 0,$$

where $\|\mathbf{x}\|_2 = \sqrt{x_1^2 + x_2^2 + \dots + x_d^2}$ for any $\mathbf{x} = (x_1, x_2, \dots, x_d) \in \mathbb{R}^d$.

Theorem 1.1. *Given any $N, L \in \mathbb{N}^+$ and a continuous function f on $[0, 1]^d$, there exists a function ϕ implemented by a Floor-ReLU network with width $\max\{d, 5N+13\}$ and depth $64dL+3$ such that*

$$|\phi(\mathbf{x}) - f(\mathbf{x})| \leq \omega_f(\sqrt{d}N^{-\sqrt{L}}) + 2\omega_f(\sqrt{d})N^{-\sqrt{L}}, \quad \text{for any } \mathbf{x} \in [0, 1]^d.$$

¹Our results can be easily generalized to Ceiling-ReLU networks, namely, feed-forward neural networks with Ceiling ($\lceil x \rceil$) and ReLU ($\max\{0, x\}$) activation functions.

The rate in $\omega_f(\sqrt{d}N^{-\sqrt{L}})$ implicitly depends on N and L through the modulus of continuity of f while the rate in $2\omega_f(\sqrt{d})N^{-\sqrt{L}}$ is explicit and independent of f . Simplifying the implicit approximation rate to make it explicitly depending on N and L is challenging in general. However, if f is a Lipschitz continuous function on $[0, 1]^d$ with a Lipschitz constant μ , then $\omega_f(r) \leq \mu r$ for any $r \geq 0$. Therefore, in the case of Lipschitz continuous functions, the approximation rate is simplified to $3\mu\sqrt{d}N^{-\sqrt{L}}$ as shown in the following corollary.

Corollary 1.2. *Given any $N, L \in \mathbb{N}^+$ and a Lipschitz continuous function f on $[0, 1]^d$ with a Lipschitz constant μ , there exists a function ϕ implemented by a Floor-ReLU network with width $\max\{d, 5N + 13\}$ and depth $64dL + 3$ such that*

$$|\phi(\mathbf{x}) - f(\mathbf{x})| \leq 3\mu\sqrt{d}N^{-\sqrt{L}}, \quad \text{for any } \mathbf{x} \in [0, 1]^d.$$

First, Theorem 1.1 and Corollary 1.2 show that the approximation capacity of deep networks for continuous functions can be exponentially improved by increasing the network depth, and the approximation error can be explicitly characterized in terms of the width $O(N)$ and depth $O(L)$. Second, this new network overcomes the curse of dimensionality in approximation power since this approximation order is essentially \sqrt{d} times a function of N and L independent of d . Finally, applying piecewise constant and integer-valued functions as activation functions and integer numbers as parameters have been explored in quantized neural networks (Hubara et al., 2017; Yin et al., 2019) with efficient training algorithms for the purpose of low computational complexity (Wang et al., 2018). As we shall see in our constructive proof of Theorem 1.1, most parameters in the Floor-ReLU network are integers. Therefore, the proposed network is also attractive for efficient computation. Though there might not be an existing optimization algorithm to identify an approximant with the approximation rate in this paper, Theorem 1.1 can provide an expected accuracy before a learning task and how much the current optimization algorithms could be improved.

Characterizing deep network approximation in terms of N and L simultaneously is fundamental and indispensable in realistic applications, while quantifying the deep network approximation based on the number of nonzero parameters W is probably only of interest in theory as far as we know. Theorem 1.1 can provide practical guidance for choosing network sizes in realistic applications while theories in terms of W cannot tell how large a network should be to guarantee a target accuracy. The width and depth are two most direct and amenable hyper-parameters in choosing a specific network for a learning task, while the number of nonzero parameters W is hardly controlled efficiently. Theories in terms of W essentially have a single variable to control the network size in three types of structures: 1) fixing the width N and varying the depth L ; 2) fixing the depth L and changing the width N ; 3) both the width and depth are controlled by the same parameter like the target accuracy ε in a specific way (e.g., N is a polynomial of $\frac{1}{\varepsilon^d}$ and L is a polynomial of $\log(\frac{1}{\varepsilon})$). Considering the non-uniqueness of structures for realizing the same W , it is impractical to develop approximation rates in terms of W covering all these structures. If one network structure has been chosen in a certain application, there might not be a known theory in terms of W to quantify the performance of this structure.

Almost all existing approximation theories for deep neural networks so far focus on the approximation rate in W (Yarotsky, 2017; Petersen and Voigtlaender, 2018; Yarotsky, 2018; Montanelli et al., 2019; Liang and Srikant, 2016; E and Wang, 2018; Opschoor et al., 2019; Barron, 1993; Montanelli and Du, 2017; Chen and Wu, 2019; Poggio et al., 2017; Yarotsky and Zhevnerchuk, 2019; Montanelli and Yang, 2020). From the point of view of theoretical difficulty, controlling two variables N and L in our theory is more challenging than controlling one variable W in the literature. In terms of mathematical logic, the characterization of deep network approximation in terms of N and L addresses the question in terms of W , while it is not true the other way around. As we have discussed in the last paragraph, existing theories essentially have a single variable to control the network size in three types of structures. Let us use the first type of structures, which includes the best-known result for a nearly optimal approximation rate $O(W^{-2/d})$ for continuous functions in terms of W (Yarotsky, 2018), as an example to show how Theorem 1.1 in terms of N and L can be applied to show a significantly much better result in terms of W . It is similar to apply Theorem 1.1 to obtain other corollaries with other types of structures in terms of W . The main idea is to specify the value of N and L in Theorem 1.1 to show the desired corollary. For example, we let the width parameter $N = 2$ and the depth parameter $L = W$ in Theorem 1.1, then the width is $\max\{d, 23\}$, the depth is $64dW + 3$, and the total number of parameters is bounded by $\max\{d^2, 23^2\}(64dW + 3) = O(W)$. Therefore, we can prove Corollary 1.3 below stating that our Floor-ReLU network can provide an approximation accuracy of $O(\sqrt{d}2^{-\sqrt{W}})$.

Corollary 1.3. *Given any $W \in \mathbb{N}^+$ and a continuous function f on $[0, 1]^d$, there exists a function ϕ implemented by a Floor-ReLU network with $O(W)$ nonzero parameters, width $\max\{d, 23\}$ and depth $64dW + 3$, such that*

$$|\phi(\mathbf{x}) - f(\mathbf{x})| \leq \omega_f(\sqrt{d}2^{-\sqrt{W}}) + 2\omega_f(\sqrt{d})2^{-\sqrt{W}}, \quad \text{for any } \mathbf{x} \in [0, 1]^d.$$

To the best of our knowledge, the neural network constructed here is the first to achieve exponential convergence and no curse of dimensionality simultaneously for a function class as general as continuous functions, while existing theories only work for functions with an intrinsic low complexity (e.g., the exponential convergence for polynomials (Yarotsky, 2017; Montanelli et al., 2019; Lu et al., 2020), smooth functions (Montanelli et al., 2019; Liang and Srikant, 2016), analytic functions (E and Wang, 2018), functions admitting a holomorphic extension to a Bernstein polyellipse (Opschoor et al., 2019); no curse of dimensionality (or the curse is lessened) for Barron spaces (Barron, 1993), Korobov spaces (Montanelli and Du, 2017), band-limited functions (Chen and Wu, 2019; Montanelli et al., 2019), compositional functions (Poggio et al., 2017), and smooth functions (Yarotsky and Zhevnerchuk, 2019; Lu et al., 2020; Montanelli and Yang, 2020; Yang and Wang, 2020)). The prefactor in our approximation rate is a known constant of size $O(\sqrt{d})$ in the case of Lipschitz continuous functions, while the prefactor of all existing theories for much smaller function classes is unknown or grows exponentially in d . Our proof fully explores the advantage of the compositional structure and the nonlinearity of deep networks, while existing theories were built on traditional approximation tools (e.g., polynomial approximation, multiresolution analysis, and Monte Carlo sampling) making it impossible for existing theories

to obtain a theoretical breakthrough. Let us review these existing works in more detail below.

In terms of no curse of dimensionality, (Barron, 1993) and its variants in (Chen and Wu, 2019; Montanelli et al., 2019) considered d -dimensional functions with Fourier integral representations, which can be approximated by the sum of N samples of the integrant at N frequencies in the same spirit of Monte Carlo sampling by the law of large numbers with an approximation error of $O(\frac{1}{\sqrt{N}})$. Target functions in (Barron, 1993; Chen and Wu, 2019; Montanelli et al., 2019) are hence required to be sufficiently smooth and the approximation error contains a prefactor that is exponentially large in d . Similarly in (Montanelli and Du, 2017), d -dimensional functions in the Korobov space are approximated by the linear combination of basis functions of a sparse grid, each of which is approximated by a ReLU network. Though the curse of dimensionality has been lessened, target functions have to be sufficiently smooth and the approximation error contains a large prefactor that is exponential in d . Similarly, the works in (Yarotsky and Zhevnerchuk, 2019; Lu et al., 2020; Yang and Wang, 2020) take advantage of the polynomial approximation to smooth functions and ReLU networks are constructed to approximate polynomials. Generally speaking, in almost all these works, the approximation power for no curse of dimensionality essentially comes from traditional tools instead of networks.

Similarly, the approximation power for exponential approximation rate in existing works comes from traditional tools for approximating a small class of functions instead of networks. In (E and Wang, 2018; Opschoor et al., 2019; Chen and Wu, 2019; Montanelli et al., 2019), highly smooth functions are first approximated by the linear combination of special polynomials with high degrees (e.g., Chebyshev polynomials, Legendre polynomials) with an exponential approximation rate, i.e., to achieve an ε -accuracy, a linear combination of only $O(p(\log(\frac{1}{\varepsilon})))$ polynomials is required, where p is a polynomial with a degree that may depend on the dimension d . Then each polynomial is approximated by a ReLU network with $O(\log(\frac{1}{\varepsilon}))$ parameters. Finally, all ReLU networks are assembled to form a large network approximating the target function with an exponential approximation rate.

Finally, deep network approximation is in fact a special case of function approximation via compositions, where an approximant space is generated as the composition of several simple latent spaces. Function compositions can significantly enhance the approximation power. The central question is to characterize the relation of the approximant space and the latent spaces so as to design simple latent spaces to generate complex approximant spaces. This is because balancing the computational complexity and the approximation capacity is crucial in designing an efficient approximation tool in realistic applications. In terms of computational efficiency, latent spaces should remain simple and structured such that efficient numerical algorithms can be designed to identify an approximant in \mathcal{S} . In terms of approximation efficiency, latent spaces should be complex enough such that they can generate a wide class of functions.

The rest of this paper is organized as follows. In Section 2, we prove Theorem 1.1 based on Proposition 2.2. Next, this basic proposition is proved in Section 3. In Section 4, we introduce function approximation via compositions as a more general framework that includes deep network approximation. Finally, we conclude this paper in Section 5.

209 2 Approximation of continuous functions

210 In this section, we first introduce basic notations in this paper in Section 2.1. Then we
 211 prove the first main theorem, Theorem 1.1, based on Proposition 2.2 in Section 3.

212 2.1 Notations

213 The main notations of this paper are listed as follows.

- 214 • Let \mathbb{N}^+ denote the set containing all positive integers, i.e., $\mathbb{N}^+ = \{1, 2, 3, \dots\}$.
- 215 • Let $\sigma : \mathbb{R} \rightarrow \mathbb{R}$ denote the rectified linear unit (ReLU), i.e. $\sigma(x) = \max\{0, x\}$.
 216 With the abuse of notations, we define $\sigma : \mathbb{R}^d \rightarrow \mathbb{R}^d$ as $\sigma(\mathbf{x}) = \begin{bmatrix} \max\{0, x_1\} \\ \vdots \\ \max\{0, x_d\} \end{bmatrix}$
 217 for any $\mathbf{x} = (x_1, \dots, x_d) \in \mathbb{R}^d$.
- 218 • For a one-dimensional function set $\Theta = \{\rho_1(x), \rho_2(x), \dots, \rho_s(x)\}$, $\boldsymbol{\varrho} \in \Theta^d$ means
 219 $\boldsymbol{\varrho}$ is a vector function of length d with each entry as a function in Θ . With the
 220 abuse of notation, $\boldsymbol{\varrho} \odot (\mathbf{x})$ is used to denote $\begin{bmatrix} \rho_1(x_1) \\ \vdots \\ \rho_d(x_d) \end{bmatrix}$ with $\rho_i \in \Theta$ for $i =$
 221 $1, \dots, d$.
- 222 • The floor function (Floor) is defined as $\lfloor x \rfloor := \max\{n : n \leq x, n \in \mathbb{Z}\}$ for any
 223 $x \in \mathbb{R}$. $\lfloor \mathbf{x} \rfloor$ means applying $\lfloor \cdot \rfloor$ entrywise to \mathbf{x} .
- 224 • For $\theta \in [0, 1)$, suppose its binary representation is $\theta = \sum_{\ell=1}^{\infty} \theta_{\ell} 2^{-\ell}$ with $\theta_{\ell} \in \{0, 1\}$,
 225 we introduce a special notation $\text{bin}0.\theta_1\theta_2\cdots\theta_L$ to denote the L -term binary repre-
 226 sentation of θ , i.e., $\sum_{\ell=1}^L \theta_{\ell} 2^{-\ell}$.
- 227 • The expression “a network ϕ ” is short of “a function ϕ that implemented by a
 228 network”.
- 229 • The expression “a network with width N and depth L ” means
 - 230 – The maximum width of this network for all hidden layers is no more than
 231 N .
 - 232 – The number of hidden layers of this network is no more than L .

233 2.2 Proof of Theorem 1.1

234 The proof of Theorem 1.1 is an immediate result of Theorem 2.1 below.

235 **Theorem 2.1.** *Given any $N, L \in \mathbb{N}^+$ and a continuous function f on $[0, 1]^d$, there exists*
 236 *a function ϕ implemented by a Floor-ReLU network with width $\max\{d, 2N^2 + 5N\}$ and*
 237 *depth $7dL^2 + 3$ such that*

$$238 \quad |\phi(\mathbf{x}) - f(\mathbf{x})| \leq \omega_f(\sqrt{d}N^{-L}) + 2\omega_f(\sqrt{d})2^{-NL}, \quad \text{for any } \mathbf{x} \in [0, 1]^d.$$

239 This theorem will be proved later in this section. Now let us prove Theorem 1.1
 240 based on Theorem 2.1.

241 *Proof of Theorem 1.1.* Given any $N, L \in \mathbb{N}^+$, there exist $\tilde{N}, \tilde{L} \in \mathbb{N}^+$ with $\tilde{N} \geq 2$ and
 242 $\tilde{L} \geq 3$ such that

$$243 \quad (\tilde{N} - 1)^2 \leq N < \tilde{N}^2 \quad \text{and} \quad (\tilde{L} - 1)^2 \leq 4L < \tilde{L}^2.$$

244 By Theorem 2.1, there exists a function ϕ implemented by a Floor-ReLU network with
 245 width $\max\{d, 2\tilde{N}^2 + 5\tilde{N}\}$ and depth $7d\tilde{L}^2 + 3$ such that

$$246 \quad |\phi(\mathbf{x}) - f(\mathbf{x})| \leq \omega_f(\sqrt{d} \tilde{N}^{-\tilde{L}}) + 2\omega_f(\sqrt{d})2^{-\tilde{N}\tilde{L}}, \quad \text{for any } \mathbf{x} \in [0, 1]^d.$$

247 Note that

$$248 \quad 2^{-\tilde{N}\tilde{L}} \leq \tilde{N}^{-\tilde{L}} = (\tilde{N}^2)^{-\frac{1}{2}\sqrt{\tilde{L}^2}} \leq N^{-\frac{1}{2}\sqrt{4L}} \leq N^{-\sqrt{L}}.$$

249 Then we have

$$250 \quad |\phi(\mathbf{x}) - f(\mathbf{x})| \leq \omega_f(\sqrt{d} N^{-\sqrt{L}}) + 2\omega_f(\sqrt{d})N^{-\sqrt{L}}, \quad \text{for any } \mathbf{x} \in [0, 1]^d.$$

251 For $\tilde{N}, \tilde{L} \in \mathbb{N}^+$ with $\tilde{N} \geq 2$ and $\tilde{L} \geq 3$, we have

$$252 \quad 2\tilde{N}^2 + \tilde{N} \leq 5(\tilde{N} - 1)^2 + 13 \leq 5N + 13 \quad \text{and} \quad 7\tilde{L}^2 \leq 16(\tilde{L} - 1)^2 \leq 64L.$$

253 Therefore, ϕ can be computed by a Floor-ReLU network with width $\max\{d, 2\tilde{N}^2 +$
 254 $5\tilde{N}\} \leq \max\{d, 5N + 13\}$ and depth $7d\tilde{L}^2 + 3 \leq 64dL + 3$, as desired. So we finish the
 255 proof. \square

256 To prove Theorem 2.1, we first present the proof sketch. Shortly speaking, we
 257 construct piecewise constant functions implemented by Floor-ReLU networks to ap-
 258 proximate continuous functions. There are six key steps in our construction.

- 259 1. Normalize f as \tilde{f} satisfying $\tilde{f}(\mathbf{x}) \in [0, 1]$ for any $\mathbf{x} \in [0, 1]^d$ and divide $[0, 1]^d$
 260 into a set of non-overlapping cubes $\{Q_\alpha\}_{\alpha \in \{0, 1, \dots, M-1\}^d}$, where M is an integer
 261 determined later.
- 262 2. Construct a vector-valued function $\Phi_1 : \mathbb{R}^d \rightarrow \mathbb{R}^d$ mapping $\mathbf{x} \in Q_\alpha$ to the index α
 263 for each $\alpha \in \{0, 1, \dots, M-1\}^d$.
- 264 3. Construct a function $\phi_2 : \mathbb{R}^d \rightarrow \mathbb{R}$ projecting $\alpha \in \{0, 1, \dots, M-1\}^d$ to $\phi_2(\alpha) \in$
 265 $\{1, 2, \dots, M^d\}$.
- 266 4. Construct a function $\phi_3 : \mathbb{R} \rightarrow \mathbb{R}$ mapping $\phi_2(\alpha) \in \{1, 2, \dots, M^d\}$ to $\phi_3(\phi_2(\alpha)) \approx$
 267 $\tilde{f}(\mathbf{x}_\alpha)$, where \mathbf{x}_α is a pre-specified point of Q_α .
- 268 5. Define $\tilde{\phi} := \phi_3 \circ \phi_2 \circ \Phi_1$. Then $\tilde{\phi}$ is a piecewise constant function mapping $\mathbf{x} \in Q_\alpha$
 269 to $\phi_3(\phi_2(\alpha)) \approx \tilde{f}(\mathbf{x}_\alpha)$.
- 270 6. Re-scale and shift $\tilde{\phi}$ to obtain the final function ϕ approximating f well.

It is not difficult to construct Floor-ReLU networks with the desired width and depth to implement Φ_1 and ϕ_2 . The most technical part is the construction of a Floor-ReLU network with the desired width and depth computing ϕ_3 , which needs the following proposition based on the “bit extraction” technique introduced in (Bartlett et al., 1998; Harvey et al., 2017).

Proposition 2.2. *Given any $N, L \in \mathbb{N}^+$ and arbitrary $\theta_m \in \{0, 1\}$ for $m = 1, 2, \dots, N^L$, there exists a function ϕ computed by a Floor-ReLU network with width $2N + 2$ and depth $7L - 2$ such that*

$$\phi(m) = \theta_m, \quad \text{for } m = 1, 2, \dots, N^L.$$

The proof of this proposition is presented in Section 3. By this proposition and the definition of VC-dimension (e.g., see (Harvey et al., 2017)), it is easy to prove that the VC-dimension of Floor-ReLU networks with constant width and depth $\mathcal{O}(L)$ has a lower bound 2^L . Such a lower bound is much larger than $\mathcal{O}(L^2)$, which is a VC-dimension upper bound of ReLU networks with the same width and depth due to Theorem 8 of (Harvey et al., 2017). This means Floor-ReLU networks are much more powerful than ReLU networks from the perspective of VC-dimension.

Based on the proof sketch stated just above, we are ready to give the detailed proof of Theorem 2.1 as follows.

Proof of Theorem 2.1. Assume f is not a constant function since it is a trivial case. Then $\omega_f(r) > 0$ for any $r > 0$. Clearly, $|f(\mathbf{x}) - f(\mathbf{0})| \leq \omega_f(\sqrt{d})$ for any $\mathbf{x} \in [0, 1]^d$. Define

$$\tilde{f} := (f - f(\mathbf{0}) + \omega_f(\sqrt{d})) / (2\omega_f(\sqrt{d})). \quad (1)$$

It follows that $\tilde{f}(\mathbf{x}) \in [0, 1]$ for any $\mathbf{x} \in [0, 1]^d$.

Set $M = N^L$, $E_{M-1} = [\frac{M-1}{M}, 1]$, and $E_m = [\frac{m}{M}, \frac{m+1}{M})$ for $m = 0, 1, \dots, M-2$. Define a step function ϕ_1 as

$$\phi_1(t) := \lfloor -\sigma(-Mt + M - 1) + M - 1 \rfloor, \quad \text{for any } t \in \mathbb{R}.^2$$

See Figure 1 for an example of ϕ_1 . It follows from the definition of ϕ_1 that

$$\phi_1(t) = m, \quad \text{if } t \in E_m, \text{ for } m = 0, 1, \dots, M-1.$$




Figure 1: An illustration of ϕ_1 on $[0, 1]$ for $M = 4$.

²If we just define $\phi_1(t) = \lfloor Mt \rfloor$, then $\phi_1(1) = M \neq M - 1$ even though $1 \in E_{M-1}$.

299 Define

$$300 \quad \Phi_1(\mathbf{x}) := (\phi_1(x_1), \phi_1(x_2), \dots, \phi_1(x_d)), \quad \text{for any } \mathbf{x} = (x_1, x_2, \dots, x_d) \in \mathbb{R}^d,$$

301 and

$$302 \quad Q_\alpha := \left\{ \mathbf{x} = (x_1, x_2, \dots, x_d) \in \mathbb{R}^d : x_j \in E_{\alpha_j} \text{ for } j = 1, 2, \dots, d \right\},$$

303 for any $\alpha = (\alpha_1, \alpha_2, \dots, \alpha_d) \in \{0, 1, \dots, M-1\}^d$. See Figure 2 for the examples of Q_α ,

304 $\alpha \in \{0, 1, \dots, M-1\}^d$, for $M = 4$ and $d = 1, 2$.

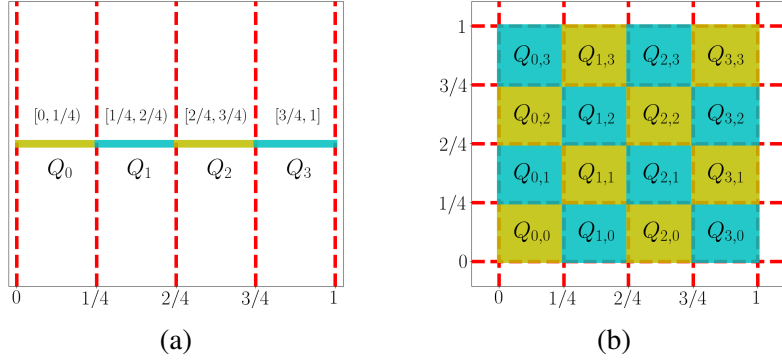


Figure 2: Illustrations of Q_α for $\alpha \in \{0, 1, \dots, M-1\}^d$. (a) $M = 4$, $d = 1$. (b) $M = 4$, $d = 2$.

305 Clearly, we have, for $\mathbf{x} \in Q_\alpha$ and $\alpha \in \{0, 1, \dots, M-1\}^d$,

$$306 \quad \Phi_1(\mathbf{x}) = (\phi_1(x_1), \phi_1(x_2), \dots, \phi_1(x_d)) = (\alpha_1, \alpha_2, \dots, \alpha_d) = \alpha.$$

307 Using the idea of M -ary representation, we define a projection function ϕ_2 via

$$308 \quad \phi_2(\mathbf{y}) := 1 + \sum_{j=1}^d y_j M^{j-1}, \quad \text{for any } \mathbf{y} = (y_1, y_2, \dots, y_d) \in \mathbb{R}^d.$$

309 Then ϕ_2 is a bijection (one-to-one correspondence) from $\{0, 1, \dots, M-1\}^d$ to $\{1, 2, \dots, M^d\}$.

310 Given any $i \in \{1, 2, \dots, M^d\}$, there exists a unique $\alpha \in \{0, 1, \dots, M-1\}^d$ such that
 311 $i = \phi_2(\alpha)$. Then define

$$312 \quad \xi_i := \tilde{f}\left(\frac{\alpha}{M}\right) \in [0, 1], \quad \text{for } i = \phi_2(\alpha) \text{ and } \alpha \in \{0, 1, \dots, M-1\}^d,$$

313 where \tilde{f} is the normalization of f defined in Equation (1). It follows that there exists
 314 $\xi_{i,\ell} \in \{0, 1\}$ for $\ell = 1, 2, \dots, NL$ such that

$$315 \quad |\xi_i - \text{bin}_{0.\xi_{i,1}\xi_{i,2}\dots\xi_{i,NL}}| \leq 2^{-NL}, \quad \text{for } i = 1, 2, \dots, M^d.$$

316 By $M^d = (N^L)^d = N^{dL}$ and Proposition 2.2, there exists a function $\phi_{3,\ell}$ computed
 317 by a Floor-ReLU network with width $2N+2$ and depth $7dL-2$, for each $\ell = 1, 2, \dots, NL$,
 318 such that

$$319 \quad \phi_{3,\ell}(i) = \xi_{i,\ell}, \quad \text{for } i = 1, 2, \dots, M^d.$$

320 By defining $\phi_3 := \sum_{\ell=1}^{NL} 2^{-\ell} \phi_{3,\ell}$, we have, for $i = \phi_2(\alpha)$ and $\alpha \in \{0, 1, \dots, M-1\}^d$,

$$\begin{aligned}
 321 \quad |\tilde{f}(\frac{\alpha}{M}) - \phi_3(\phi_2(\alpha))| &= |\xi_i - \phi_3(i)| = |\xi_i - \sum_{\ell=1}^{NL} 2^{-\ell} \phi_{3,\ell}(i)| \\
 &= |\xi_i - \text{bin}_{0.\xi_{i,1}\xi_{i,2}\dots\xi_{i,NL}}| \leq 2^{-NL}.
 \end{aligned} \tag{2}$$

322 Define $\tilde{\phi} := \phi_3 \circ \phi_2 \circ \Phi_1$, i.e., for any $\mathbf{x} = (x_1, x_2, \dots, x_d) \in \mathbb{R}^d$,

$$323 \quad \tilde{\phi}(\mathbf{x}) = \phi_3 \circ \phi_2 \circ \Phi_1(\mathbf{x}) = \phi_3\left(\phi_2(\phi_1(x_1), \phi_1(x_2), \dots, \phi_1(x_d))\right).$$

324 Note that $|\mathbf{x} - \frac{\alpha}{M}| \leq \frac{\sqrt{d}}{M}$ for any $\mathbf{x} \in Q_\alpha$ and $\alpha \in \{0, 1, \dots, M-1\}^d$. Then we have,
 325 for any $\mathbf{x} \in Q_\alpha$ and $\alpha \in \{0, 1, \dots, M-1\}^d$,

$$\begin{aligned}
 |\tilde{f}(\mathbf{x}) - \tilde{\phi}(\mathbf{x})| &\leq |\tilde{f}(\mathbf{x}) - \tilde{f}(\frac{\alpha}{M})| + |\tilde{f}(\frac{\alpha}{M}) - \tilde{\phi}(\mathbf{x})| \\
 &\leq \omega_{\tilde{f}}(\frac{\sqrt{d}}{M}) + |\tilde{f}(\frac{\alpha}{M}) - \phi_3(\phi_2(\Phi_1(\mathbf{x})))| \\
 326 \quad &\leq \omega_{\tilde{f}}(\frac{\sqrt{d}}{M}) + |\tilde{f}(\frac{\alpha}{M}) - \phi_3(\phi_2(\alpha))| \leq \omega_{\tilde{f}}(\frac{\sqrt{d}}{M}) + 2^{-NL},
 \end{aligned}$$

327 where the last inequality comes from Equation (2).

328 Note $\mathbf{x} \in Q_\alpha$ and $\alpha \in \{0, 1, \dots, M-1\}^d$ are arbitrary. Since $[0, 1]^d = \cup_{\alpha \in \{0, 1, \dots, M-1\}^d} Q_\alpha$,
 329 we have

$$330 \quad |\tilde{f}(\mathbf{x}) - \tilde{\phi}(\mathbf{x})| \leq \omega_{\tilde{f}}(\frac{\sqrt{d}}{M}) + 2^{-NL}, \quad \text{for any } \mathbf{x} \in [0, 1]^d.$$

331 Define $\phi := 2\omega_f(\sqrt{d})\tilde{\phi} + f(\mathbf{0}) - \omega_f(\sqrt{d})$. By $M = N^L$ and $\omega_f(r) = 2\omega_f(\sqrt{d}) \cdot \omega_{\tilde{f}}(r)$
 332 for any $r \geq 0$, we have, for any $\mathbf{x} \in [0, 1]^d$,

$$\begin{aligned}
 |f(\mathbf{x}) - \phi(\mathbf{x})| &= 2\omega_f(\sqrt{d})|\tilde{f}(\mathbf{x}) - \tilde{\phi}(\mathbf{x})| \leq 2\omega_f(\sqrt{d})\left(\omega_{\tilde{f}}(\frac{\sqrt{d}}{M}) + 2^{-NL}\right) \\
 333 \quad &\leq \omega_f(\frac{\sqrt{d}}{M}) + 2\omega_f(\sqrt{d})2^{-NL} \\
 &\leq \omega_f(\sqrt{d}N^{-L}) + 2\omega_f(\sqrt{d})2^{-NL}.
 \end{aligned}$$

334 Since ϕ is defined via re-scaling and shifting $\tilde{\phi}$, what remains is to determine the
 335 width and depth of the Floor-ReLU network computing $\tilde{\phi}$. Clearly, ϕ_3 can be implemented by the architecture in Figure 3.




Figure 3: An illustration of the desired network architecture computing ϕ_3 . We omit some ReLU (σ) activation functions if inputs are obviously non-negative.




Figure 4: An illustration of the desired network architecture computing $\tilde{\phi}$, where $\tilde{\phi}(\mathbf{x}) = \phi_3 \circ \phi_2 \circ \Phi_1(\mathbf{x})$ for any $\mathbf{x} = (x_1, x_2, \dots, x_d) \in \mathbb{R}^d$. We omit some ReLU (σ) activation functions if inputs are obviously non-negative.

As we can see from Figure 3, ϕ_3 can be computed by a Floor-ReLU network with width $N(2N + 2 + 3) = 2N^2 + 5N$ and depth $L(7dL - 2 + 1) + 1 = L(7dL - 1) + 1$. With the network architecture computing ϕ_3 in hand, $\tilde{\phi}$ can be implemented by the network architecture shown in Figure 4.

By Figure 4, ϕ and $\tilde{\phi}$ can be implemented by a Floor-ReLU network with width $\max\{d, 2N^2 + 5N\}$ and depth $L(7dL - 1) + 1 + 3 \leq 7dL^2 + 3$. So we finish the proof. \square

3 Proof of Proposition 2.2

The proof of Proposition 2.2 mainly relies on the “bit extraction” technique. As we shall see later, our key idea is to apply the Floor activation function to make “bit extraction” more powerful to reduce network sizes. In particular, Floor-ReLU networks can extract much more bits than ReLU networks with the same network size.

Let us first establish a basic lemma to extract $1/N$ of total bits stored in a new binary number from an input binary number.

Lemma 3.1. *Given any $J, N \in \mathbb{N}^+$, there exists a function $\phi : \mathbb{R}^2 \rightarrow \mathbb{R}$ that can be implemented by a Floor-ReLU network with width $2N$ and depth 4 such that, for any $\theta_j \in \{0, 1\}$, $j = 1, \dots, NJ$, we have*

$$\phi(\text{bin}0.\theta_1 \dots \theta_{NJ}, n) = \text{bin}0.\theta_{(n-1)J+1} \dots \theta_{nJ}, \quad \text{for } n = 1, 2, \dots, N.$$

Proof. Given any $\theta_j \in \{0, 1\}$ for $j = 1, \dots, NJ$, denote

$$s = \text{bin}0.\theta_1 \dots \theta_{NJ} \quad \text{and} \quad s_n = \text{bin}0.\theta_{(n-1)J+1} \dots \theta_{nJ}, \quad \text{for } n = 1, 2, \dots, N.$$

Then our goal is to construct a function $\phi : \mathbb{R}^2 \rightarrow \mathbb{R}$ computed by a Floor-ReLU network with the desired width and depth that satisfies

$$\phi(s, n) = s_n, \quad \text{for } n = 1, 2, \dots, N.$$

Based on the properties of the binary representation, it is easy to check that

$$s_n = \lfloor 2^{nJ} s \rfloor / 2^{(n-1)J} - \lfloor 2^{(n-1)J} s \rfloor, \quad \text{for } n = 1, 2, \dots, N. \quad (3)$$

With formulas to return s_1, s_2, \dots, s_N , it is still technical to construct a network outputting s_n for a given index $n \in \{1, 2, \dots, N\}$.




Figure 5: An illustration of $g(x) = \sigma\left(\sigma(x) - \sigma\left(\frac{x+\delta-1}{\delta}\right)\right)$, where $\sigma(x) = \max\{0, x\}$.

363 Set $\delta = 2^{-NJ}$ and define g (see Figure 5) as

364
$$g(x) := \sigma\left(\sigma(x) - \sigma\left(\frac{x+\delta-1}{\delta}\right)\right), \quad \text{where } \sigma(x) = \max\{0, x\}.$$

365 Since $s_n \in [0, 1 - \delta]$ for $n = 1, 2, \dots, N$, we have

366
$$s_n = \sum_{k=1}^N g(s_k + k - n), \quad \text{for } n = 1, 2, \dots, N. \quad (4)$$




Figure 6: This architecture is based on Equation (3) and (4). We omit some ReLU (σ) activation functions provided inputs are obviously non-negative.

367 As shown in Figure 6, the desired function ϕ can be computed by a Floor-ReLU
368 network with width $2N$ and depth 4. Moreover, it holds that

369
$$\phi(s, n) = s_n, \quad \text{for } n = 1, 2, \dots, N.$$

370 So we finish the proof. □

371 The next lemma constructs a Floor-ReLU network that can extract any bit from a
372 binary number according to a specific index.

373 **Lemma 3.2.** *Given any $N, L \in \mathbb{N}^+$, there exists a function $\phi : \mathbb{R}^2 \rightarrow \mathbb{R}$ implemented by a*
374 *Floor-ReLU network with width $2N + 2$ and depth $7L - 3$ such that, for any $\theta_m \in \{0, 1\}$,*
375 *$m = 1, 2, \dots, N^L$, we have*

376
$$\phi(\text{bin}0.\theta_1\theta_2\cdots\theta_{N^L}, m) = \theta_m, \quad \text{for } m = 1, 2, \dots, N^L.$$

377 *Proof.* The proof is based on repeated applications of Lemma 3.1. To be exact, we
378 construct a sequence of functions $\phi_1, \phi_2, \dots, \phi_L$ implemented by Floor-ReLU networks
379 by induction to satisfy the following two conditions for each $\ell \in \{1, 2, \dots, L\}$.

380 (i) $\phi_\ell : \mathbb{R}^2 \rightarrow \mathbb{R}$ can be implemented by a Floor-ReLU network with width $2N + 2$
 381 and depth $7\ell - 3$.

382 (ii) For any $\theta_m \in \{0, 1\}$, $m = 1, 2, \dots, N^\ell$, we have

$$383 \quad \phi_\ell(\text{bin}0.\theta_1\theta_2\cdots\theta_{N^\ell}, m) = \text{bin}0.\theta_m, \quad \text{for } m = 1, 2, \dots, N^\ell.$$

384 Firstly, consider the case $\ell = 1$. By Lemma 3.1 (set $J = 1$ therein), there exists a
 385 function ϕ_1 implemented by a Floor-ReLU network with width $4 \leq 6$ and depth $4 = 7 - 3$
 386 such that, for any $\theta_m \in \{0, 1\}$, $m = 1, 2, \dots, N$, we have

$$387 \quad \phi_1(\text{bin}0.\theta_1\theta_2\cdots\theta_N, m) = \text{bin}0.\theta_m, \quad \text{for } m = 1, 2, \dots, N.$$

388 It follows that Condition (i) and (ii) hold for $\ell = 1$.

389 Next, assume Condition (i) and (ii) hold for $\ell = k$. We would like to construct ϕ_{k+1}
 390 to make Condition (i) and (ii) true for $\ell = k + 1$. By Lemma 3.1 (set $J = N^k$ therein),
 391 there exists a function ψ implemented by a Floor-ReLU network with width $2N + 2$ and
 392 depth 4 such that, for any $\theta_m \in \{0, 1\}$, $m = 1, 2, \dots, N^k$, we have

$$393 \quad \psi(\text{bin}0.\theta_1\cdots\theta_{N^{k+1}}, n) = \text{bin}0.\theta_{(n-1)N^k+1}\cdots\theta_{(n-1)N^k+N^k}, \quad \text{for } n = 1, 2, \dots, N. \quad (5)$$

394 By the hypothesis of induction, we have

- 395 • $\phi_k : \mathbb{R}^2 \rightarrow \mathbb{R}$ can be implemented by a Floor-ReLU network with width $2N + 2$
 396 and depth $7k - 3$.
- 397 • For any $\theta_j \in \{0, 1\}$, $j = 1, 2, \dots, N^k$, we have

$$398 \quad \phi_k(\text{bin}0.\theta_1\theta_2\cdots\theta_{N^k}, j) = \text{bin}0.\theta_j, \quad \text{for } j = 1, 2, \dots, N^k. \quad (6)$$

399 Given any $m \in \{1, 2, \dots, N^{k+1}\}$, there exist $n \in \{1, 2, \dots, N\}$ and $j \in \{1, 2, \dots, N^k\}$
 400 such that $m = (n - 1)N^k + j$, and such n, j can be obtained by

$$401 \quad n = \lfloor (m - 1)/N^k \rfloor + 1 \quad \text{and} \quad j = m - (n - 1)N^k. \quad (7)$$

402 Then the desired architecture of the Floor-ReLU network implementing ϕ_{k+1} is shown
 in Figure 7.




Figure 7: This architecture is based on Equation (5), (6), and (7). We omit some ReLU (σ) activation functions provided inputs are obviously non-negative.

403 Note that ψ can be computed by a Floor-ReLU network with width $2N$ and depth
 404 4. By Figure 7, we have
 405

406 • $\phi_{k+1} : \mathbb{R}^2 \rightarrow \mathbb{R}$ can be implemented by a Floor-ReLU network with width $2N + 2$
 407 and depth $7 + (7k - 3) = 7(k + 1) - 3$, which implies Condition (i) for $\ell = k + 1$.

408 • For any $\theta_m \in \{0, 1\}$, $m = 1, 2, \dots, N^{k+1}$, we have

$$409 \quad \phi_{k+1}(\text{bin}0.\theta_1\theta_2\cdots\theta_{N^{k+1}}, m) = \text{bin}0.\theta_m, \quad \text{for } m = 1, 2, \dots, N^{k+1}.$$

410 That is, Condition (ii) holds for $\ell = k + 1$.

411 So we finish the process of induction.

412 By the principle of induction, there exists a function $\phi_L : \mathbb{R}^2 \rightarrow \mathbb{R}$ such that

413 • ϕ_L can be implemented by a Floor-ReLU network with width $2N + 2$ and depth
 414 $7L - 3$.

415 • For any $\theta_m \in \{0, 1\}$, $m = 1, 2, \dots, N^L$, we have

$$416 \quad \phi_L(\text{bin}0.\theta_1\theta_2\cdots\theta_{N^L}, m) = \text{bin}0.\theta_m, \quad \text{for } m = 1, 2, \dots, N^L.$$

417 Finally, define $\phi := 2\phi_L$. Then ϕ can also be implemented by a Floor-ReLU network
 418 with width $2N + 2$ and depth $7L - 3$. Moreover, for any $\theta_m \in \{0, 1\}$, $m = 1, 2, \dots, N^L$,
 419 we have

$$420 \quad \phi(\text{bin}0.\theta_1\theta_2\cdots\theta_{N^L}, m) = 2 \cdot \phi_L(\text{bin}0.\theta_1\theta_2\cdots\theta_{N^L}, m) = 2 \cdot \text{bin}0.\theta_m = \theta_m,$$

421 for $m = 1, 2, \dots, N^L$. So we finish the proof. \square

422 With Lemma 3.2 in hand, we are ready to prove Proposition 2.2.

423 *Proof of Proposition 2.2.* By Lemma 3.2, there exists a function $\tilde{\phi} : \mathbb{R}^2 \rightarrow \mathbb{R}$ computed
 424 by a Floor-ReLU network with a fixed architecture with width $2N + 2$ and depth $7L - 3$
 425 such that, for any $z_m \in \{0, 1\}$, $m = 1, 2, \dots, N^L$, we have

$$426 \quad \tilde{\phi}(\text{bin}0.z_1z_2\cdots z_{N^L}, m) = z_m, \quad \text{for } m = 1, 2, \dots, N^L.$$

427 Based on $\theta_m \in \{0, 1\}$ for $m = 1, 2, \dots, N^L$ given in Proposition 2.2, we define the final
 428 function ϕ as

$$429 \quad \phi(x) := \tilde{\phi}(\sigma(x \cdot 0 + \text{bin}0.\theta_1\theta_2\cdots\theta_{N^L}), \sigma(x)), \quad \text{where } \sigma(x) = \max\{0, x\}.$$

430 Clearly, ϕ can be implemented by a Floor-ReLU network with width $2N + 2$ and depth
 431 $(7L - 3) + 1 = 7L - 2$. Moreover, we have, for any $m \in \{1, 2, \dots, N^L\}$,

$$432 \quad \phi(m) := \tilde{\phi}(\sigma(m \cdot 0 + \text{bin}0.\theta_1\theta_2\cdots\theta_{N^L}), \sigma(m)) = \tilde{\phi}(\text{bin}0.\theta_1\theta_2\cdots\theta_{N^L}, m) = \theta_m.$$

433 So we finish the proof. \square

434 We shall point out that only the properties of Floor on $[0, \infty)$ are used in our proof.
 435 Thus, the Floor can be replaced by the truncation function that can be easily computed
 436 by truncating the decimal part.

4 Approximation via compositions

In this section, we will discuss function compositions for approximation and its a few examples including deep network approximation. Let us first formulate the problem of function approximation via compositions, where an approximant space \mathcal{S} is generated by the composition of latent spaces \mathcal{S}_ℓ as follows.

Definition 4.1. Suppose $N_0 = d$, $N_{L+1} = k$, $N_\ell \in \mathbb{N}_+$, and \mathcal{S}_ℓ is a space of functions from $\mathbb{R}^{N_{\ell-1}}$ to \mathbb{R}^{N_ℓ} for $\ell = 1, \dots, L+1$. Let

$$\mathcal{S} := \{\mathbf{h}(\mathbf{x}) = \mathbf{h}_{L+1} \circ \mathbf{h}_L \circ \dots \circ \mathbf{h}_1(\mathbf{x}) : \mathbf{h}_\ell \in \mathcal{S}_\ell \text{ for } \ell = 1, \dots, L+1\},$$

then \mathcal{S} is called the approximant space generated by latent spaces \mathcal{S}_ℓ for $\ell = 1, \dots, L+1$ with L compositions.

A best approximant of $\mathbf{f}(\mathbf{x})$ in \mathcal{S} is identified by solving

$$\phi(\mathbf{x}) = \arg \min_{\phi(\mathbf{x}) \in \mathcal{S}} \|\mathbf{f}(\mathbf{x}) - \phi(\mathbf{x})\|_*, \quad (8)$$

where $\|\cdot\|_*$ is an appropriate norm depending on applications. Function compositions can significantly enhance the approximation power. Composing a fixed latent space several times could generate a much richer approximant space. However, this idea was not considered in the literature previously due to the expensive computation in solving the optimization problem (8). Deep learning and its related optimization algorithms (e.g., backpropagation techniques (Werbos, 1975; Fukushima, 1980; Rumelhart et al., 1986), parallel computing techniques (Scherer et al., 2010; Cireřan et al., 2011), and stochastic algorithms (Duchi et al., 2011; Johnson and Zhang, 2013)) indicate that solving Equation (8) has become feasible and hence compositions can be a practical choice for function approximation.

4.1 Classical approximation

Though function approximation via compositions is relatively new, most existing approximation techniques can be considered as its special cases.

Linear approximation

Let us first discuss the linear approximation through the lens of approximation via compositions. Linear approximation is an efficient approximation tool for smooth functions that computes the approximant of a target function via a linear projection to a Hilbert space or a Banach space as the approximant space. The linear projection can be computed efficiently through the orthogonality of basis functions in the Hilbert space or via (quasi) interpolation in the Banach space. Typically, target functions are required to be sufficiently smooth to obtain a good numerical approximation via the projection to a fixed set of finitely many (e.g., N) basis functions. Approximation theories and numerical tools have been well-developed in linear approximation, e.g. polynomial approximations, Fourier analysis, finite element approximation, spline approximation, etc. They have become powerful tools for approximating smooth functions.

Linear approximation can be considered as a special case of function approximation via compositions when $L = 1$, $k = 1$, and $N_1 = N$. For simplicity, let us take the one-dimensional orthogonal polynomial approximation as an example. The approximant space \mathcal{S} , in this case, is generated by

$$\mathcal{S}_1 = \{\boldsymbol{\varrho}(x) = [\varrho_0(x), \varrho_1(x), \dots, \varrho_{N-1}(x)]^T\},$$

and

$$\mathcal{S}_2 = \{f(\mathbf{x}) = \mathbf{W} \cdot \mathbf{x} : \mathbf{W} \in \mathbb{R}^N\}, \quad (9)$$

where $\{\varrho_i(x)\}_{i=0}^{N-1}$ is a set of orthogonal polynomials with degrees from zero to $N - 1$. The computation between a target function and an approximant is highly efficient due to the orthogonality.

Similarly, when the target function space is $L^2(\mathbb{R}^d)$ and \mathcal{S}_1 consists of a vector function of length N with entries as basis functions in the Fourier series of \mathbb{R}^d . N here restricts approximants to the first N -term Fourier series expansion. Fourier basis functions admit a good structure, orthogonality, which makes it simple to compute an approximant in \mathcal{S} to approximate a target function f and to reconstruct f from its representation.

In general, the computational complexity of tools in the case of $L = 1$ could be as low as nearly optimal, e.g., $O(N)$ operations ignoring a logarithm factor. Nevertheless, their approximation capacity suffers from the curse of dimensionality and there is no exponential approximation rate for general continuous functions. The approximation accuracy of \mathcal{S} when $L = 1$ and $N_1 = N$ is usually $O(N^{-1/d})$ for a continuous function, which is far from $O(N^{-\sqrt{L}})$ when d is large.

Nonlinear approximation

Nonlinear approximation (DeVore, 1998) has become a popular technique in recent decades for piecewise-smooth function approximation. A typical algorithm in nonlinear approximation is to design a highly redundant nonlinear dictionary, \mathcal{D} , and to identify the optimal approximant as a linear combination of N elements of \mathcal{D} . Given a dictionary \mathcal{D} and a target function $f(\mathbf{x})$, nonlinear approximation seeks $\{g_n\}$ and $\{T_n\}$ such that

$$\{\{T_n\}, \{g_n\}\} = \arg \min_{\{g_n\} \subseteq \mathbb{R}, \{T_n\} \subseteq \mathcal{D}} \|f(\mathbf{x}) - \sum_{n=1}^N g_n T_n(\mathbf{x})\|_*, \quad (10)$$

which is also called the best N -term approximation with an appropriate norm $\|\cdot\|_*$.

Traditional dictionaries in nonlinear approximation can be considered as a special case of function approximation via compositions when $L = 1$, $k = 1$, and $N_1 = N$. Wavelet frames of $L^2([0, 1]^d)$ built with the dilation and translation of a mother wavelet can serve as a typical example. The approximant space \mathcal{S} of wavelet frames can be generated by latent spaces

$$\mathcal{S}_1 = \{\mathbf{h}(\mathbf{x}) = \boldsymbol{\varrho} \odot (\mathbf{W} \cdot \mathbf{x} + \mathbf{b}) : \mathbf{W} \in \mathbb{R}^{N \times d}, \mathbf{b} \in \mathbb{R}^N, \boldsymbol{\varrho} \in \Theta^N, 1 \leq i \leq N\},$$

and

$$\mathcal{S}_2 = \{h(\mathbf{x}) = \mathbf{W} \cdot \mathbf{x} : \mathbf{W} \in \mathbb{R}^N\},$$

where $\Theta = \{\varrho(x)\}$ with $\varrho(x)$ as a mother wavelet, \mathbf{W} corresponds to dilation, and \mathbf{b} determines translation. Though wavelet frames are redundant and hence there is no orthogonality, under the assumption that target functions have a sparse approximant, it is

still computationally efficient to determine a best approximant in \mathcal{S} and reconstruct a target function via the ℓ_1 -regularization to Equation (10) in realistic applications. Similar to the case of linear approximation, the approximation capacity of traditional nonlinear approximation suffers from the curse of dimensionality and there is no exponential approximation rate for general continuous functions, e.g. typically $O(N^{-1/d})$ for a continuous function.

Nonlinear approximation via compositions proposed in (Shen et al., 2019) can provide a more attractive approximation rate. The key idea is to use function compositions to generate a dictionary in nonlinear approximation. For example, in the case when $N_L = N$, $k = 1$, and \mathcal{S}_{L+1} is defined by Equation (9), seeking a best function approximation via compositions by solving Equation (8) is equivalent to Equation (10) when we let

$$\mathcal{D} = \mathcal{D}_L := \{[\mathbf{h}_L \circ \cdots \circ \mathbf{h}_1(\mathbf{x})]_j : \mathbf{h}_\ell \in \mathcal{S}_\ell, \ell = 1, \dots, L, j = 1, \dots, N\},$$

where $[\mathbf{h}(\mathbf{x})]_j$ means the j -th output of the vector function $\mathbf{h}(\mathbf{x})$. Nonlinear approximation concerns the quantification of the best N -term approximation rate in N defined as

$$\varepsilon_f(N) = \min_{\{g_n\} \subseteq \mathbb{R}, \{T_n\} \subseteq \mathcal{D}} \|f(\mathbf{x}) - \sum_{n=1}^N g_n T_n(\mathbf{x})\|_*.$$

Hence, when we use \mathcal{D}_L as the dictionary, it is interesting to quantify

$$\varepsilon_{L,f}(N) = \min_{\{g_n\} \subseteq \mathbb{R}, \{T_n\} \subseteq \mathcal{D}_L} \|f(\mathbf{x}) - \sum_{n=1}^N g_n T_n(\mathbf{x})\|_*.$$

Function compositions can significantly enrich the dictionary of nonlinear approximation. For example, if the dictionary is built with the Floor-ReLU networks proposed in this paper, $\varepsilon_{L,f}(N) \leq O(N^{-\sqrt{L}})$ for Lipschitz continuous functions f . This rate is much better than existing rates for much smaller function classes, e.g. $O(N^{-s/d})$ for functions in Besov spaces with smoothness s (DeVore and Ron, 2010; Hangelbroek and Ron, 2010), and $O(N^{-\frac{1}{2d}})$ for Hölder continuous functions of order 1 on $[0, 1]^d$ (Xie and Cao, 2013).

4.2 Approximation by compositions

Since approximation via compositions has not been fully explored yet, there are many new research directions remaining. The central question is to characterize the relation of the approximant space \mathcal{S} and the latent spaces \mathcal{S}_ℓ so as to guide the design of latent spaces according to the requirement of the approximant space. In terms of computational efficiency, latent spaces \mathcal{S}_ℓ should remain simple and structured such that it is easy to parametrize \mathcal{S}_ℓ with an efficient numerical algorithm to identify an approximant in \mathcal{S} . In terms of approximation efficiency, latent spaces should be complex enough such that they can generate a wide class of functions.

Generating an approximant space with a large number of latent spaces is a natural preference balancing the computational complexity and the approximation capacity. There is a dilemma making it difficult to design a powerful approximant space \mathcal{S} when L is small. For example, when $L = 1$, it is required that \mathcal{S}_1 is sufficiently simple and complex simultaneously. However, when L is large, there is much room to use simple

latent spaces \mathcal{S}_ℓ to generate a complicated approximant space \mathcal{S} , which is the most prevailing advantage of function approximation via compositions.

For example, the approximant space of deep ReLU networks with depth L can be generated by latent spaces defined as

$$\mathcal{S}_\ell = \{h(\mathbf{x}) = \boldsymbol{\varrho} \odot (\mathbf{W} \cdot \mathbf{x} + \mathbf{b}) : \mathbf{W} \in \mathbb{R}^{N_\ell \times N_{\ell-1}}, \mathbf{b} \in \mathbb{R}^{N_\ell}, \boldsymbol{\varrho} \in \Theta^{N_\ell}\}, \quad (11)$$

with $\Theta = \{\max\{0, x\}\}$, $1 \leq i \leq N_\ell$, $\ell = 1, \dots, L$, and

$$\mathcal{S}_{L+1} = \{h(\mathbf{x}) = \mathbf{W} \cdot \mathbf{x} + \mathbf{b} : \mathbf{W} \in \mathbb{R}^{k \times N_L}, \mathbf{b} \in \mathbb{R}^k\}.$$

Recall that, for a set of one-dimensional functions Θ , $\boldsymbol{\varrho} \odot (\mathbf{x}) \in \Theta^d$ is used to denote

$$\boldsymbol{\varrho} \odot (\mathbf{x}) = \begin{bmatrix} \rho_1(x_1) \\ \vdots \\ \rho_d(x_d) \end{bmatrix} \text{ with } \rho_i \in \Theta \text{ for } i = 1, \dots, d.$$

Similarly, the space of Floor-ReLU neural networks studied in this paper can be generated by latent spaces \mathcal{S}_ℓ defined as

$$\mathcal{S}_\ell = \{h(\mathbf{x}) = \boldsymbol{\varrho} \odot (\mathbf{W} \cdot \mathbf{x} + \mathbf{b}) : \mathbf{W} \in \mathbb{R}^{N_\ell \times N_{\ell-1}}, \mathbf{b} \in \mathbb{R}^{N_\ell}, \boldsymbol{\varrho} \in \Theta^{N_\ell}\}, \quad (12)$$

with $\Theta = \{\max\{0, x\}, \lfloor x \rfloor\}$, $1 \leq i \leq N_\ell$, for $\ell = 1, \dots, L$, and

$$\mathcal{S}_{L+1} = \{h(\mathbf{x}) = \mathbf{W} \cdot \mathbf{x} + \mathbf{b} : \mathbf{W} \in \mathbb{R}^{k \times N_L}, \mathbf{b} \in \mathbb{R}^k\}.$$

Motivated by Floor-ReLU networks, it would be interesting to investigate the approximant space $\mathcal{S}_{\Theta, N}$ generated by latent spaces \mathcal{S}_ℓ defined via

$$\mathcal{S}_\ell = \{h(\mathbf{x}) = \boldsymbol{\varrho} \odot (\mathbf{W} \cdot \mathbf{x} + \mathbf{b}) : \mathbf{W} \in \mathbb{R}^{N_\ell \times N_{\ell-1}}, \mathbf{b} \in \mathbb{R}^{N_\ell}, \boldsymbol{\varrho} \in \Theta^{N_\ell}\} \quad (13)$$

for $\ell = 1, \dots, L$, and

$$\mathcal{S}_{L+1} = \{h(\mathbf{x}) = \mathbf{W} \cdot \mathbf{x} + \mathbf{b} : \mathbf{W} \in \mathbb{R}^{k \times N_L}, \mathbf{b} \in \mathbb{R}^k\},$$

where $\mathbf{N} = (N_0, N_1, \dots, N_L, N_{L+1}) \in \mathbb{N}_+^{L+2}$, $N_0 = d$, $N_{L+1} = k$, and Θ is a finite set of one-dimensional functions. The formulation in Equation (13) keeps it neat and easy for computation. Though functions in \mathcal{S}_ℓ defined in Equation (13) are high-dimensional, they only consist of a simple high-dimensional linear transform and a few one-dimensional nonlinear functions. Therefore, the essential complexity of each function in \mathcal{S}_ℓ is $O(N_\ell)$. The formulation in Equation (13) is sufficiently powerful to make it unnecessary to consider more complex latent spaces, because any continuous latent space can be approximated by the compositions of latent spaces in Equation (13) as shown by Theorem 1.1.

There are mainly two research directions to characterize the approximation power of compositions: 1) fixing one kind of latent spaces and varying target function spaces; 2) fixing a target function space as general as possible and explore different kinds of latent spaces. For simplicity, let us assume $N_\ell = N$ for $\ell = 1, \dots, L$ and $k = 1$ in Definition 4.1.

In the first research direction, there has been extensive research in the literature starting from shallow neural networks with the Sigmoid activation function to deep neural

networks with the ReLU activation function (Cybenko, 1989; Hornik et al., 1989; Barron, 1993; Yarotsky, 2018, 2017; Bölcskei et al., 2019; Zhou, 2019; Chui et al., 2018; Gribonval et al., 2019; Gühring et al., 2019; Suzuki, 2019; Nakada and Imaizumi, 2019; Chen et al., 2019; Bao et al., 2019; Li et al., 2019; Montanelli and Yang, 2020). For example, it was shown in (Lu et al., 2020) that deep ReLU neural networks can achieve an approximation rate $O(N^{-L})$ for multi-dimensional polynomials, but $O(N^{-L})$ is not true for general smooth functions, e.g. a nearly optimal approximation rate for $C^s([0, 1]^d)$ functions is $O(N^{-2s/d})$. For a larger function class, e.g., $C([0, 1]^d)$ functions, a nearly optimal approximation rate of ReLU networks is $O(N^{-2/d})$ as shown in (Shen et al., 2019). Hence, the first research direction with an existing network structure may not provide an appealing approximation rate unless the target function space is sufficiently small and structured.

The second research direction is more promising and open to new theories. It was stated in (Yarotsky and Zhevnerchuk, 2019) with a sketchy discussion that, when Θ in Equation (13) contains both the ReLU and sin functions, the approximant space $\mathcal{S}_{\Theta, \mathbf{N}}$ with $\max\{N_1, \dots, N_L\} = O(1)$ can approximate smooth functions on $[0, 1]^d$ with an approximation rate $O(e^{-\sqrt{L}})$. Though the power of function compositions can be reflected by the exponent \sqrt{L} , the contribution of width is missing and the exponent is not proportional to L . Furthermore, the target function space is much smaller than the continuous function space. When Θ in Equation (13) contains both the ReLU and Floor activation functions, by applying Theorem 1.1 for k times and assembling the resulting k networks into a large network, we can show that an approximation rate of $O(N^{-\sqrt{L}})$ is achievable as in Corollary 4.2 below. It would be very interesting to explore other Θ 's to see whether $O(N^{-L})$ is achievable and what is the key characterization of latent spaces to achieve this rate.

Corollary 4.2. *There exist $\Theta = \{\max\{0, x\}, \lfloor x \rfloor\}$ such that given any $N, L \in \mathbb{N}^+$, for an arbitrary continuous vector function $\mathbf{f}(\mathbf{x}) \in \mathbb{R}^k$ on $[0, 1]^d$, there exists a function $\phi \in \mathcal{S}_{\Theta, \mathbf{N}}$ with $\mathbf{N} = \{d, N_1, \dots, N_{\tilde{L}}, k\}$, $\tilde{L} = 64dL + 3$, and $\max\{dk, k(5N + 13)\} \geq \max\{N_1, \dots, N_{\tilde{L}}\}$ such that*

$$\|\phi(\mathbf{x}) - \mathbf{f}(\mathbf{x})\|_{\ell^\infty} \leq \omega_{\mathbf{f}}(\sqrt{d}N^{-\sqrt{L}}) + 2\omega_{\mathbf{f}}(\sqrt{d})N^{-\sqrt{L}}$$

for any $\mathbf{x} \in [0, 1]^d$, where $\omega_{\mathbf{f}}(x) := \max\{\omega_{f_1}(x), \dots, \omega_{f_k}(x)\}$.

5 Conclusion

This paper has introduced the first theoretical framework to show that deep network approximation can achieve exponential convergence and avoid the curse of dimensionality for approximating functions as general as continuous functions. Given a Lipschitz continuous function f on $[0, 1]^d$, it was shown by construction that Floor-ReLU networks with width $\max\{d, 5N + 13\}$ and depth $64dL + 3$ admit a uniform approximation rate $3\mu\sqrt{d}N^{-\sqrt{L}}$, where μ is the Lipschitz constant of f . More generally for an arbitrary continuous function f on $[0, 1]^d$ with a modulus of continuity $\omega_f(\cdot)$, the constructive approximation rate is $\omega_f(\sqrt{d}N^{-\sqrt{L}}) + 2\omega_f(\sqrt{d})N^{-\sqrt{L}}$. Function approximation via

630 compositions was also introduced as a more general framework including deep network
631 approximation.

632 **Acknowledgments.** Z. Shen is supported by Tan Chin Tuan Centennial Professor-
633 ship. H. Yang was partially supported by the US National Science Foundation under
634 award DMS-1945029.

635 References

- 636 Bao, C., Li, Q., Shen, Z., Tai, C., Wu, L., and Xiang, X. (2019). Approximation analysis
637 of convolutional neural networks.
- 638 Barron, A. R. (1993). Universal approximation bounds for superpositions of a sigmoidal
639 function. *IEEE Transactions on Information Theory*, 39(3):930–945.
- 640 Bartlett, P., Maierov, V., and Meir, R. (1998). Almost linear VC-dimension bounds for
641 piecewise polynomial networks. *Neural Computation*, 10:217–3.
- 642 Bölcskei, H., Grohs, P., Kutyniok, G., and Petersen, P. (2019). Optimal approximation
643 with sparsely connected deep neural networks. *SIAM Journal on Mathematics of*
644 *Data Science*, 1(1):8–45.
- 645 Chen, L. and Wu, C. (2019). A note on the expressive power of deep rectified linear
646 unit networks in high-dimensional spaces. *Mathematical Methods in the Applied*
647 *Sciences*, 42(9):3400–3404.
- 648 Chen, M., Jiang, H., Liao, W., and Zhao, T. (2019). Efficient approximation of deep relu
649 networks for functions on low dimensional manifolds. In Wallach, H., Larochelle,
650 H., Beygelzimer, A., d'Alché-Buc, F., Fox, E., and Garnett, R., editors, *Advances*
651 *in Neural Information Processing Systems 32*, pages 8174–8184. Curran Associates,
652 Inc.
- 653 Chui, C. K., Lin, S.-B., and Zhou, D.-X. (2018). Construction of neural networks
654 for realization of localized deep learning. *Frontiers in Applied Mathematics and*
655 *Statistics*, 4:14.
- 656 Cireşan, D. C., Meier, U., Masci, J., Gambardella, L. M., and Schmidhuber, J. (2011).
657 Flexible, high performance convolutional neural networks for image classification.
658 In *Proceedings of the Twenty-Second International Joint Conference on Artificial*
659 *Intelligence - Volume Volume Two, IJCAI'11*, pages 1237–1242. AAAI Press.
- 660 Cybenko, G. (1989). Approximation by superpositions of a sigmoidal function. *MCSS*,
661 2:303–314.
- 662 DeVore, R. and Ron, A. (2010). Approximation using scattered shifts of a multivariate
663 function. *Transactions of the American Mathematical Society*, 362(12):6205–6229.
- 664 DeVore, R. A. (1998). Nonlinear approximation. *Acta Numerica*, 7:51–150.

665 Duchi, J., Hazan, E., and Singer, Y. (2011). Adaptive subgradient methods for online
666 learning and stochastic optimization. *J. Mach. Learn. Res.*, 12:2121–2159.

667 E, W. and Wang, Q. (2018). Exponential convergence of the deep neural network ap-
668 proximation for analytic functions. *CoRR*, abs/1807.00297.

669 Fukushima, K. (1980). Neocognitron: A self-organizing neural network model for a
670 mechanism of pattern recognition unaffected by shift in position. *Biological Cyber-*
671 *netics*, 36(4):193–202.

672 Gribonval, R., Kutyniok, G., Nielsen, M., and Voigtlaender, F. (2019). Approximation
673 spaces of deep neural networks. *arXiv e-prints*, page arXiv:1905.01208.

674 Gühring, I., Kutyniok, G., and Petersen, P. (2019). Error bounds for approxi-
675 mations with deep ReLU neural networks in $W^{s,p}$ norms. *arXiv e-prints*, page
676 arXiv:1902.07896.

677 Hangelbroek, T. and Ron, A. (2010). Nonlinear approximation using gaussian kernels.
678 *Journal of Functional Analysis*, 259(1):203 – 219.

679 Harvey, N., Liaw, C., and Mehrabian, A. (2017). Nearly-tight VC-dimension bounds
680 for piecewise linear neural networks. In Kale, S. and Shamir, O., editors, *Proceedings*
681 *of the 2017 Conference on Learning Theory*, volume 65 of *Proceedings of Machine*
682 *Learning Research*, pages 1064–1068, Amsterdam, Netherlands. PMLR.

683 Hornik, K., Stinchcombe, M., and White, H. (1989). Multilayer feedforward networks
684 are universal approximators. *Neural Networks*, 2(5):359 – 366.

685 Hubara, I., Courbariaux, M., Soudry, D., El-Yaniv, R., and Bengio, Y. (2017). Quan-
686 tized neural networks: Training neural networks with low precision weights and ac-
687 tivations. *J. Mach. Learn. Res.*, 18(1):6869–6898.

688 Johnson, R. and Zhang, T. (2013). Accelerating stochastic gradient descent using pre-
689 dictive variance reduction. In *Proceedings of the 26th International Conference on*
690 *Neural Information Processing Systems - Volume 1*, NIPS’13, pages 315–323, USA.
691 Curran Associates Inc.

692 Li, Q., Lin, T., and Shen, Z. (2019). Deep Learning via Dynamical Systems: An Ap-
693 proximation Perspective. *arXiv e-prints*, page arXiv:1912.10382.

694 Liang, S. and Srikant, R. (2016). Why deep neural networks? *CoRR*, abs/1610.04161.

695 Lu, J., Shen, Z., Yang, H., and Zhang, S. (2020). Deep Network Approximation for
696 Smooth Functions. *arXiv e-prints*, page arXiv:2001.03040.

697 Montanelli, H. and Du, Q. (2017). New error bounds for deep networks using sparse
698 grids.

699 Montanelli, H. and Yang, H. (2020). Error bounds for deep relu networks using the
700 kolmogorov–arnold superposition theorem. *Neural Networks*, 129:1 – 6.

- 701 Montanelli, H., Yang, H., and Du, Q. (2019). Deep ReLU networks overcome the curse
702 of dimensionality for bandlimited functions.
- 703 Nakada, R. and Imaizumi, M. (2019). Adaptive approximation and estimation of deep
704 neural network with intrinsic dimensionality.
- 705 Opschoor, J. A., Schwab, C., and Zech, J. (2019). Exponential ReLU DNN expression
706 of holomorphic maps in high dimension. Technical report, Zurich.
- 707 Petersen, P. and Voigtlaender, F. (2018). Optimal approximation of piecewise smooth
708 functions using deep ReLU neural networks. *Neural Networks*, 108:296 – 330.
- 709 Poggio, T., Mhaskar, H. N., Rosasco, L., Miranda, B., and Liao, Q. (2017). Why and
710 when can deep—but not shallow—networks avoid the curse of dimensionality: A
711 review. *International Journal of Automation and Computing*, 14:503–519.
- 712 Rumelhart, D., McClelland, J., Group, P. R., and University of California, S. D. P. R. G.
713 (1986). *Psychological and Biological Models*. Number v. 2 in A Bradford Book.
714 MIT Press.
- 715 Scherer, D., Müller, A., and Behnke, S. (2010). Evaluation of pooling operations in
716 convolutional architectures for object recognition. In Diamantaras, K., Duch, W.,
717 and Iliadis, L. S., editors, *Artificial Neural Networks – ICANN 2010*, pages 92–101,
718 Berlin, Heidelberg. Springer Berlin Heidelberg.
- 719 Shen, Z., Yang, H., and Zhang, S. (2019). Deep network approximation characterized
720 by number of neurons. *arXiv e-prints*, page arXiv:1906.05497.
- 721 Shen, Z., Yang, H., and Zhang, S. (2019). Nonlinear approximation via compositions.
722 *Neural Networks*, 119:74 – 84.
- 723 Suzuki, T. (2019). Adaptivity of deep reLU network for learning in besov and mixed
724 smooth besov spaces: optimal rate and curse of dimensionality. In *International*
725 *Conference on Learning Representations*.
- 726 Wang, P., Hu, Q., Zhang, Y., Zhang, C., Liu, Y., and Cheng, J. (2018). Two-step quan-
727 tization for low-bit neural networks. In *2018 IEEE/CVF Conference on Computer*
728 *Vision and Pattern Recognition*, pages 4376–4384.
- 729 Werbos, P. (1975). *Beyond Regression: New Tools for Prediction and Analysis in the*
730 *Behavioral Sciences*. Harvard University.
- 731 Xie, T. F. and Cao, F. L. (2013). The rate of approximation of gaussian radial basis
732 neural networks in continuous function space. *Acta Mathematica Sinica, English*
733 *Series*, 29(2):295–302.
- 734 Yang, Y. and Wang, Y. (2020). Approximation in shift-invariant spaces with deep ReLU
735 neural networks. *arXiv e-prints*, page arXiv:2005.11949.
- 736 Yarotsky, D. (2017). Error bounds for approximations with deep ReLU networks. *Neu-*
737 *ral Networks*, 94:103 – 114.

- 738 Yarotsky, D. (2018). Optimal approximation of continuous functions by very deep
739 ReLU networks. In Bubeck, S., Perchet, V., and Rigollet, P., editors, *Proceedings*
740 *of the 31st Conference On Learning Theory*, volume 75 of *Proceedings of Machine*
741 *Learning Research*, pages 639–649. PMLR.
- 742 Yarotsky, D. and Zhevnerchuk, A. (2019). The phase diagram of approximation rates
743 for deep neural networks. *arXiv e-prints*, page arXiv:1906.09477.
- 744 Yin, P., Lyu, J., Zhang, S., Osher, S., Qi, Y., and Xin, J. (2019). Understand-
745 ing straight-through estimator in training activation quantized neural nets. *ArXiv*,
746 abs/1903.05662.
- 747 Zhou, D.-X. (2019). Universality of deep convolutional neural networks. *Applied and*
748 *Computational Harmonic Analysis*.